

Saehee Eom

Atlanta, GA • (404) 203-5710 • saehee.data@gmail.com • <https://saeheeeom.github.io/>



EDUCATION

GEORGIA INSTITUTE OF TECHNOLOGY (GPA: 4.0/4.0)

Atlanta, GA

Master of Science in Computer Science (Machine Learning)

Aug. 2025 - Dec. 2026

Relevant Coursework: Machine Learning, Language Models & Agents, Statistical Modeling, Database Systems

Graduate TA, CSE 6242 Data and Visual Analytics

SEOUL NATIONAL UNIVERSITY (GPA: 3.7/4.0, Cum Laude)

Seoul, South Korea

Bachelor of Science in Artificial Intelligence, Bachelor of Arts in Linguistics

Mar. 2019 - Feb. 2025

Relevant Coursework: AI, Natural Language Processing, Computer Vision, Machine Learning, Database

EXPERIENCE

GT Human-Centered AI Lab, Research Intern

Feb. 2026 - Present

- Built a stratified **~5B-token pretraining corpus** (576 topic × format bins) to test whether attributed training data **causally drives LLM behavior**; Co-authored a **COLM 2026** paper
- Implemented **machine unlearning (NGDiff)** from paper descriptions and ran large-scale unlearning experiments on **OLMo-3 7B** across **Slurm HPC clusters**

GE Vernova, Data Scientist Intern

Jun. 2026 - Aug. 2026

- Designed a rubric-based evaluation framework for production LLM agents, decomposing system prompts into atomic criteria and scoring **1,750 evaluation points** to raise instruction adherence from **~50% to 97–98%**
- Created AI agents and automation workflows (n8n) that parse contract and pricing data from scattered PDFs into renewal reports, cutting manual processing time by **~80%** across **20+ contracts**

Salesforce, Data Scientist Intern

Jan. 2025 - Jun. 2025

- Developed a probabilistic ML scoring model on historical opportunity data to predict deal close likelihood by sales stage and industry segment, ranking high-conversion opportunities
- Constructed bottom-up TAM/SAM models across **50+ industries**; analysis adopted by management for next FY sales planning and headcount allocation

LBOX, Data Scientist Intern

Jul. 2022 - Jan. 2023

- Co-authored an NLLP@EMNLP 2022 paper on information extraction for legal documents, building baseline models and evaluation frameworks

PROJECTS

Cross-Lingual LLM Negotiation | NeurIPS 2026 SocialAgent Workshop (First Author)

- Designed a factorial experiment running **860+** multi-turn conversational LLM negotiations via Google Vertex AI and OpenRouter across two model families
- Engineered an LLM-as-judge pipeline annotating **430+** transcripts, detecting a systematic mislabeling bias and fixing it through iterative prompt revision
- Built a parallelized, resume-safe experiment orchestrator with hardened API retry and error handling, **recovering 91 failed runs** without full reruns

Personalized Recommendation System | Hyundai Home Shopping Networks

- Boosted target-user purchase rate by **42%** via a LightGBM + SASRec ensemble deployed on AWS; translated model insights into targeted campaigns, contributing to a **14%** lift in campaign engagement

Hybrid Vector Search & RAG System | Georgia Tech, Database Systems

- Designed a LangChain RAG pipeline over **150K+** documents combining FAISS vector search with SQL filtering, achieving **3–30×** lower latency via a learned query router (**~96% F1**)

PUBLICATIONS

Same Persona, Different Language, Different Deal: Cross-Lingual Effects in LLM Negotiation, SocialAgent@NeurIPS'26

Where Does Social Reasoning Come From? Capability Provenance in Language Models, COLM'26 [[arXiv](#)]

Data-efficient End-to-end Information Extraction for Statistical Legal Analysis, NLLP@EMNLP'22 [[arXiv](#)]

SKILLS

ML & Modeling: Python (Pandas, Scikit-Learn, PyTorch), SQL, Recommendation & Ranking, Predictive Modeling

GenAI & Agents: AI Agents & Workflow Automation (n8n), LLM Evaluation (LLM-as-Judge, Rubric-based), Prompt Engineering, RAG (LangChain, FAISS), LLM APIs (Vertex AI, OpenRouter, OpenAI)

Tools: HuggingFace, Slurm, Snowflake, Spark (PySpark), AWS, GCP, Git, Tableau